



AI VIET NAM

@aivietnam.edu.vn

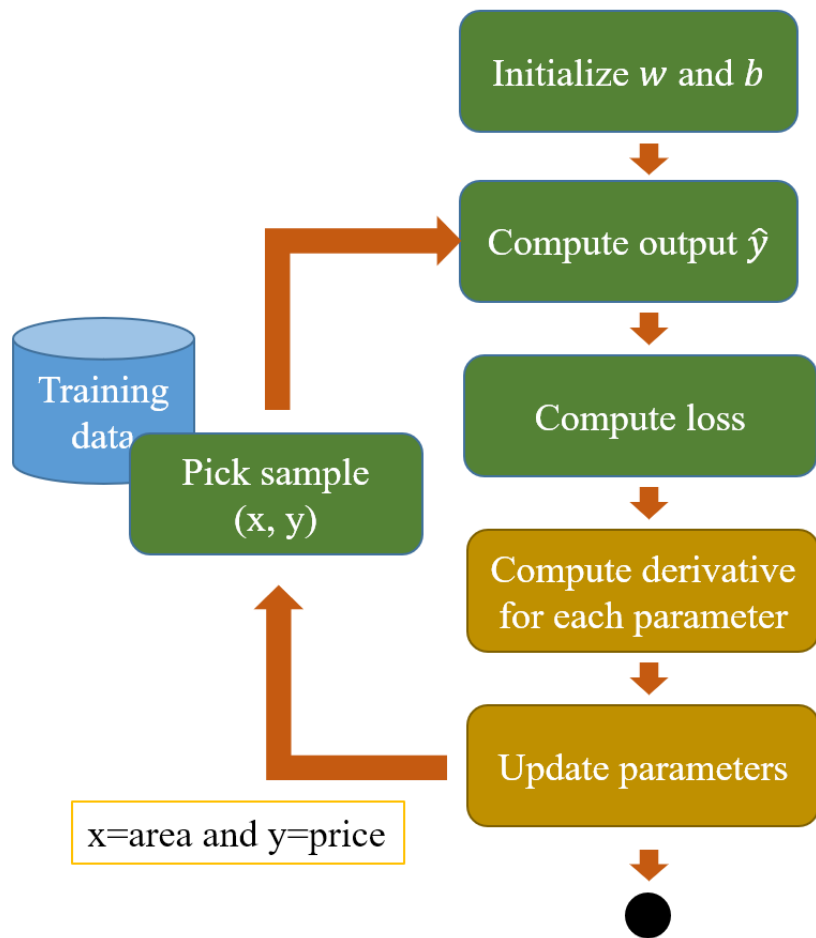
Linear Regression

A Simple Approach


Quang-Vinh Dinh
PhD in Computer Science

Objectives

Linear Regression



Computational Graph



Batch Training

1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < N$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$
$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)}) \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

Outline

SECTION 1

Linear Regression

SECTION 2

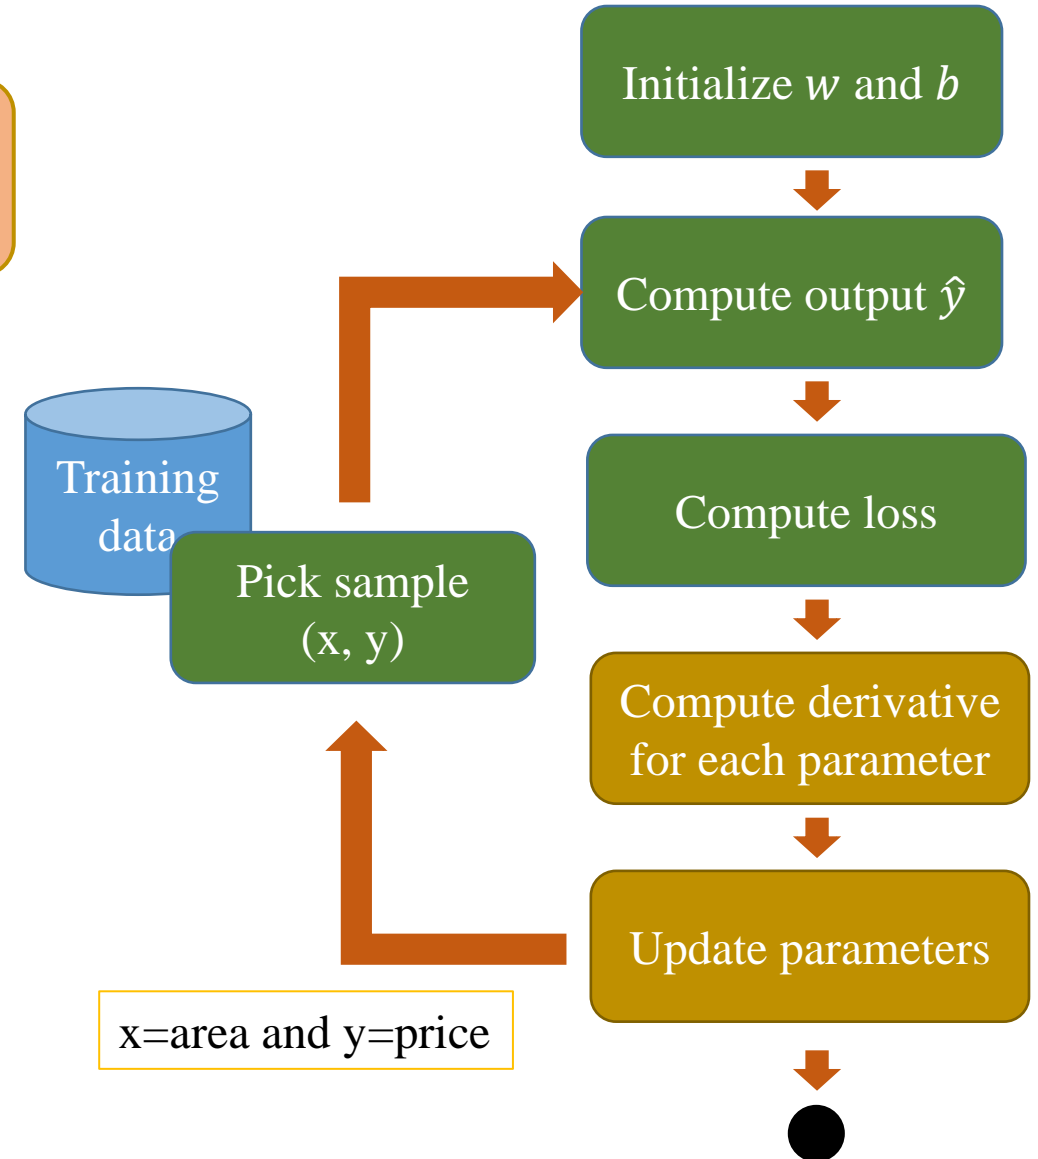
Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions



Linear Regression

Introduction

| Feature | | Label | |
|---------|------|-------|--|
| | area | price | |
| | 6.7 | 9.1 | |
| | 4.6 | 5.9 | |
| | 3.5 | 4.6 | |
| | 5.5 | 6.7 | |
| | | | |

House price data

| Features | | | Label |
|----------|-------|-----------|-------|
| TV | Radio | Newspaper | Sales |
| 230.1 | 37.8 | 69.2 | 22.1 |
| 44.5 | 39.3 | 45.1 | 10.4 |
| 17.2 | 45.9 | 69.3 | 12 |
| 151.5 | 41.3 | 58.5 | 16.5 |
| 180.8 | 10.8 | 58.4 | 17.9 |

Advertising data

if area=6.0, price=?

if TV=55.0, Radio=34.0,
and Newspaper=62.0,
price=?

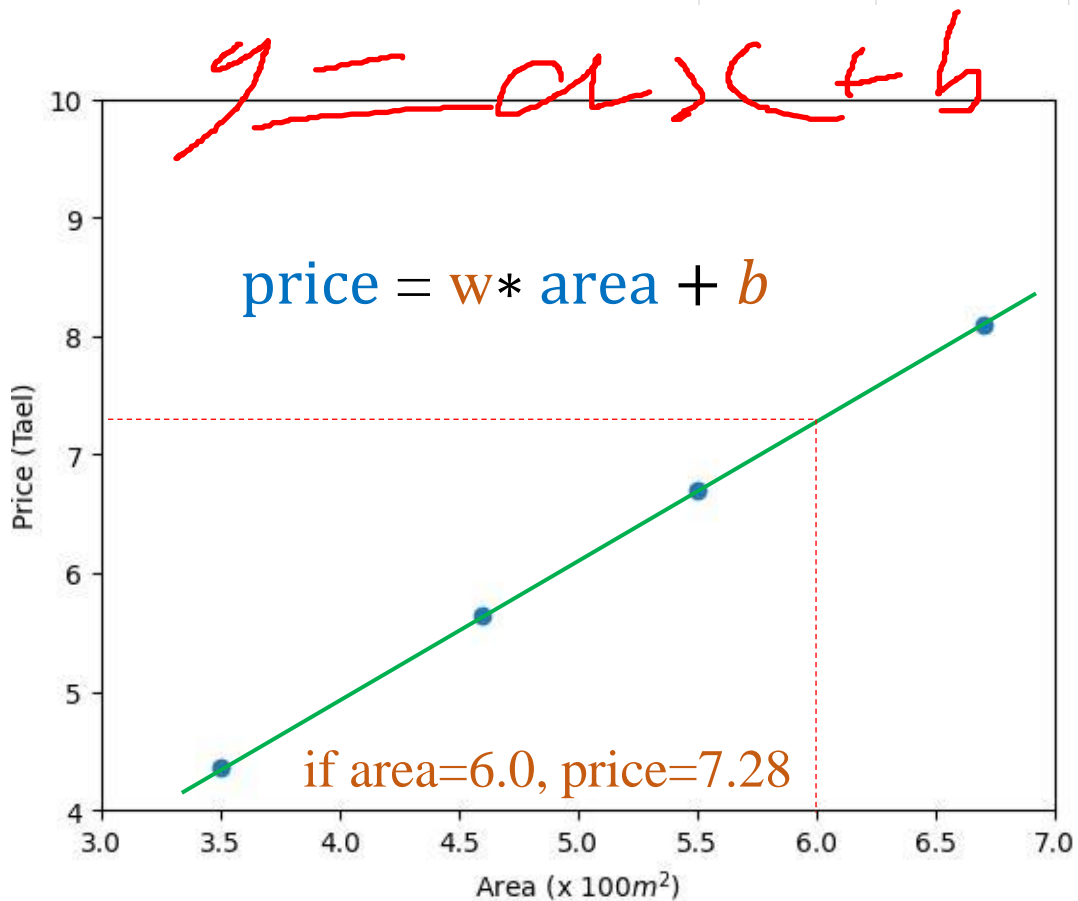
| Features | | | | | | | | | | | | | Label |
|----------|------|-------|------|-------|-------|------|--------|-----|-----|---------|--------|-------|-------|
| crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | black | lstat | medv |
| 0.00632 | 18 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.09 | 1 | 296 | 15.3 | 396.9 | 4.98 | 24 |
| 0.02731 | 0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 396.9 | 9.14 | 21.6 |
| 0.03237 | 0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 394.63 | 2.94 | 33.4 |
| 0.06905 | 0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 396.9 | 5.33 | 36.2 |
| 0.08829 | 12.5 | 7.87 | 0 | 0.524 | 6.012 | 66.6 | 5.5605 | 5 | 311 | 15.2 | 395.6 | 12.43 | 22.9 |

Boston House Price Data

House Price Prediction


| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 8.1 |
| 4.6 | 5.6 |
| 3.5 | 4.3 |
| 5.5 | 6.7 |

House price data



| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |

House price data






Linear Regression

3

❖ Area-based house price prediction



$error_i = distance(\hat{y}_i, y_i)$



Find w and b whose model has the smallest error, where $error = \frac{1}{N} \sum_i error_i$ **How?**

❖ Area-based house price prediction

weight

bias

predicted_price = w * area + b


error = (predicted_price - real_price)²

$\hat{y}_i = wx_i + b$

$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$

| area | price | predicted | error |
|------|-------|-----------|--------|
| 6.7 | 9.1 | -2.238 | 128.55 |
| 4.6 | 5.9 | -1.524 | 55.11 |
| 3.5 | 4.6 | -1.15 | 33.06 |
| 5.5 | 6.7 | -1.83 | 72.76 |

$w = -0.34$
 $b = 0.04$



❖ Area-based house price prediction


$$\text{predicted_price} = w * \text{area} + b$$

$$\text{error} = (\text{predicted_price} - \text{real_price})^2$$

$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$



| area | price | predicted | error |
|------|-------|-----------|---------|
| 6.7 | 9.1 | 8.099 | 1.002 |
| 4.6 | 5.9 | 5.642 | 0.066 |
| 3.5 | 4.6 | 4.355 | 0.06 |
| 5.5 | 6.7 | 6.695 | 0.00002 |



❖ Area-based house price prediction

$$\hat{y}_i = wx_i + b$$
$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b
so that $L(\hat{y}_i, y_i)$ reduces



Linear Regression


$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$


❖ Understanding the loss function

w *b*

How to change w and b so that $L(\hat{y}_i, y_i)$ reduces



Different b values with a fixed w value



Different w values with a fixed b value

Linear equation

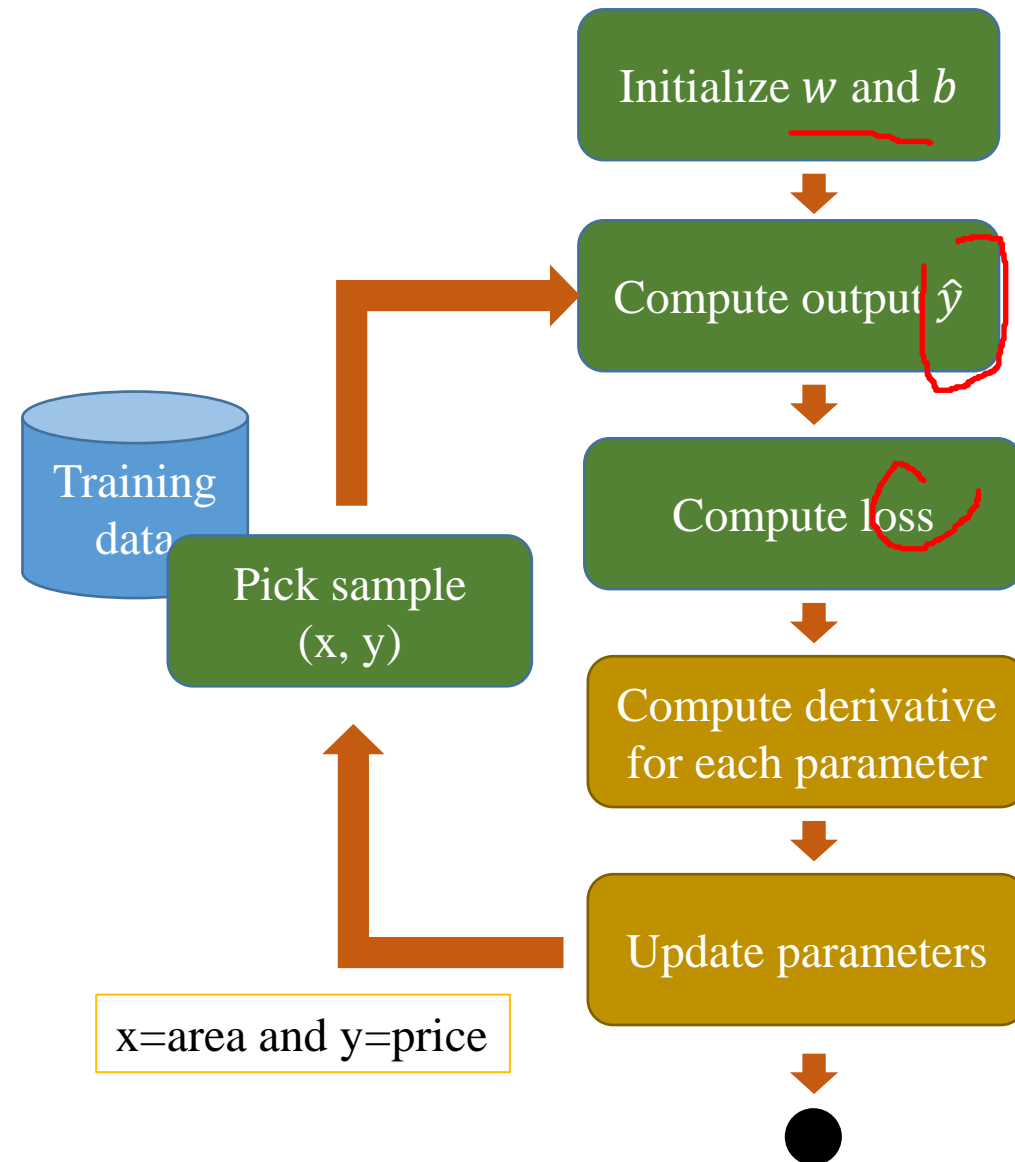
$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,
 w and b are parameters
and x is an input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y
Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$



Linear Regression

Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

and x is an input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

Find better w and b

Use gradient descent to minimize the loss function

Compute derivate for each parameter

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

Update parameters

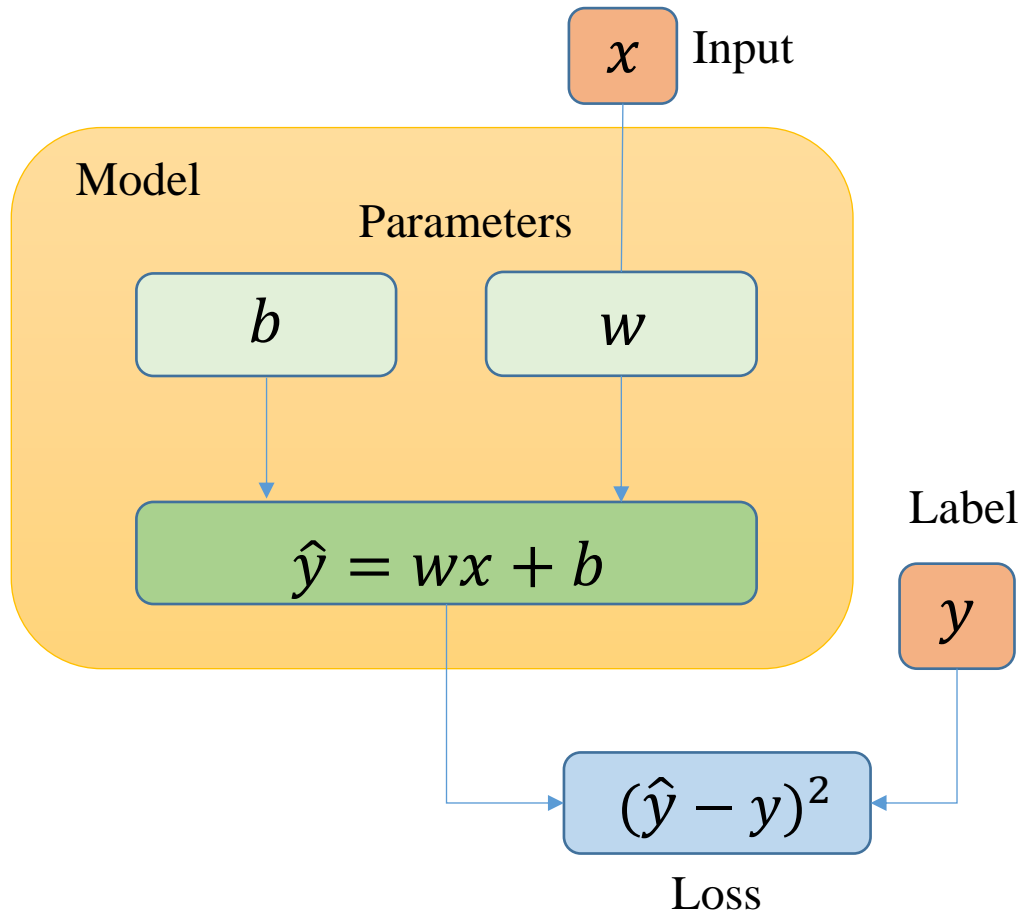
$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

❖ Simple example

Diagram



Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$


$$b = b - \eta \frac{\partial L}{\partial b}$$

Linear Regression

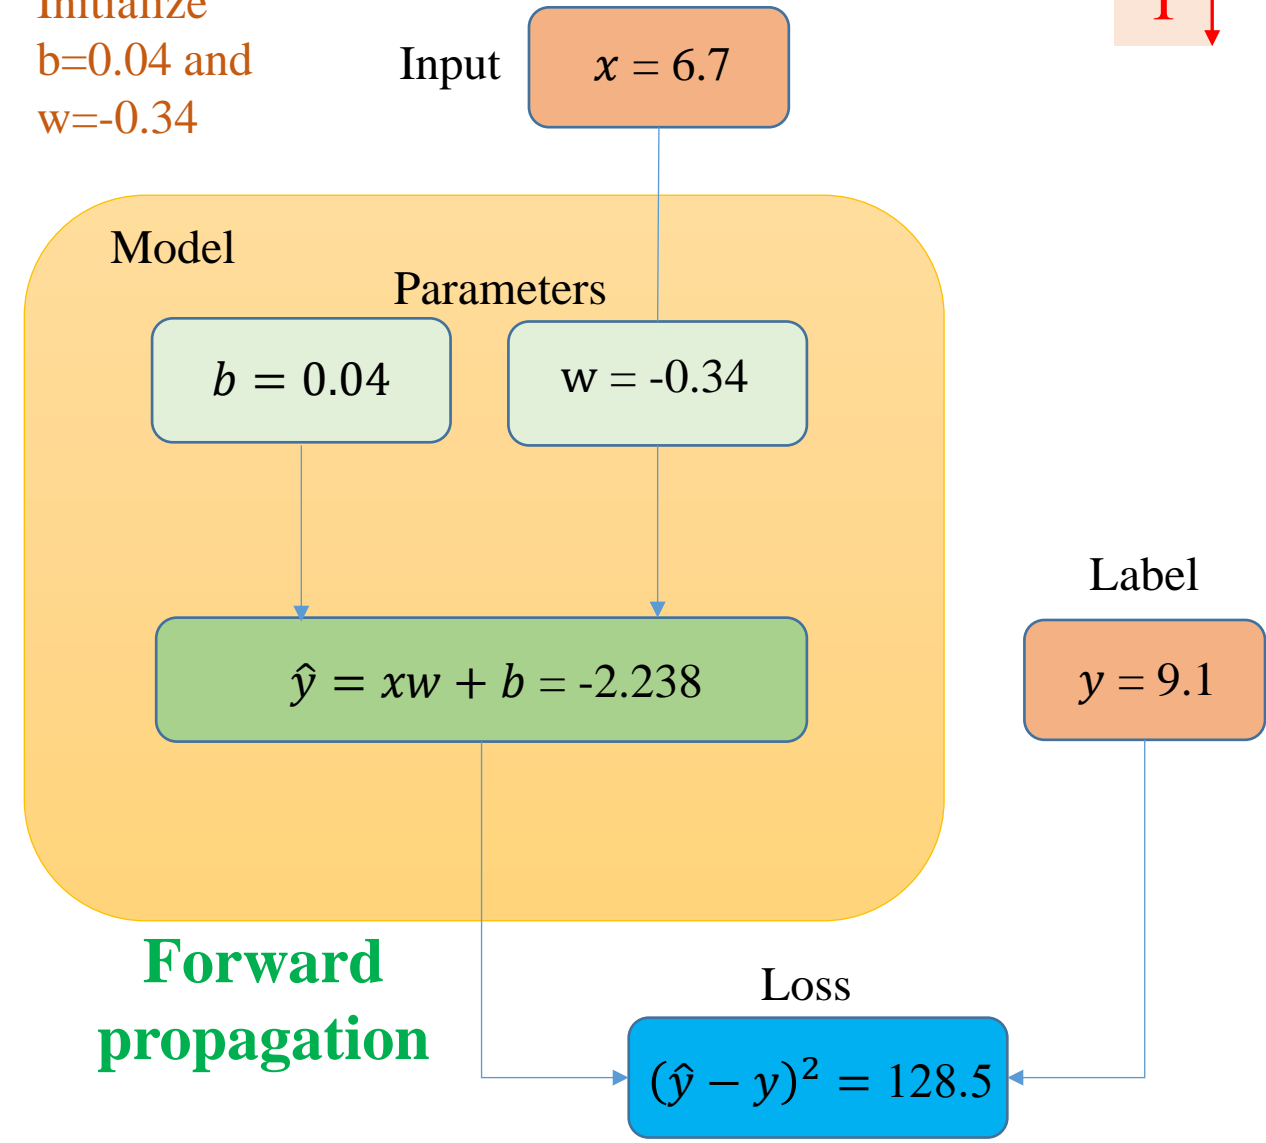
11

Given
sample
data

| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |



Initialize
 $b=0.04$ and
 $w=-0.34$



Linear Regression

12

2

Input

$x = 6.7$

Backpropagation

$\eta = 0.01$

Model

Parameters

$b = 0.26676$

$w = 1.17929$

$$b = b - \eta \frac{\partial L}{\partial b}$$

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$\hat{y} = xw + b = -2.238$$

$$\begin{aligned} \frac{\partial L}{\partial w} &= 2x(\hat{y} - y) \\ &= -151.9292 \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial b} &= 2(\hat{y} - y) \\ &= -22.676 \end{aligned}$$

Label

$y = 9.1$

Loss

$$(\hat{y} - y)^2 = 128.5$$

3

Input

$x = 6.7$

Forward propagation

Model

Parameters

$b = 0.26676$

$w = 1.17929$

$$b = b - \eta \frac{\partial L}{\partial b}$$

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$\hat{y} = xw + b = 8.168$$

Label

$y = 9.1$


Loss

$$(\hat{y} - y)^2 = 0.868$$

New w and b help
the loss reduce


❖ Simple example

Model prediction before and after the first update



$w = -0.34$ $b = 0.04$ $L = 128.55$


Before updating



$w = 1.179292$ $b = 0.26676$ $L = 0.868$

After updating

❖ Summary (simple version)



1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

❖ Implementation

Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

```
1 # forward
2 def predict(x, w, b):
3     return x*w + b
4
5 # compute gradient
6 def gradient(y_hat, y, x):
7     dw = 2*x*(y_hat-y)
8     db = 2*(y_hat-y)
9
10    return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17    return (w_new, b_new)
```


Computational Graph (A different viewpoint)

❖ House price prediction

❖ One-sample training

$$price = w * area + b$$

$$t = w * area$$



Computational graph

17

❖ House price prediction


❖ One-sample training

$$\hat{y} = wx + b$$

$$t = wx$$

$$\hat{y} = t + b$$

$$L = (\hat{y} - y)^2$$



Initial values

 $b = 0.04$ $w = -0.34$ $t = wx$ $\hat{y} = t + b$ $L = (\hat{y} - y)^2$ dL


$$\frac{dw}{dy} = 2(\hat{y} - y)$$

w and b are usually set by random values

For example, $N(0, \sigma)$ where σ is small

❖ House price prediction


❖ One-sample training



| | Feature | Label | |
|--|---------|-------|--|
| | area | price | |
| | 6.7 | 9.1 | |
| | 4.6 | 5.9 | |
| | 3.5 | 4.6 | |
| | 5.5 | 6.7 | |
| | | | |


❖ House price prediction

❖ One-sample training



❖ House price prediction

❖ One-sample training



Update w and b

$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$


Learning rate $\eta = 0.01$

$$w = -0.34 - (0.01 * (-151.9)) = 1.179$$

$$b = 0.04 - (0.01 * (-22.67)) = 0.266$$

❖ House price prediction

❖ One-sample training




| | Feature | Label | |
|--|---------|-------|--|
| | area | price | |
| | 6.7 | 9.1 | |
| | 4.6 | 5.9 | |
| | 3.5 | 4.6 | |
| | 5.5 | 6.7 | |
| | | | |

❖ House price prediction

❖ One-sample training

| | Feature | Label |
|--|---------|-------|
| | area | price |
| | 6.7 | 9.1 |
| | 4.6 | 5.9 |
| | 3.5 | 4.6 |
| | 5.5 | 6.7 |



Updated a and b values help to reduce the L value

Implementation

❖ One-sample training

| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |
| | |
| | |

column index=0 column index=1

```
1  # data preparation
2  import numpy as np
3  import matplotlib.pyplot as plt
4
5  def get_column(data, index):
6      result = [row[index] for row in data]
7      return result
8
9  data = np.genfromtxt('data.csv',
10                      delimiter=',').tolist()
11
12  x_data = get_column(data, 0)
13  y_data = get_column(data, 1)
14  N = len(x_data)
15
16  print(f'areas: {x_data}')
17  print(f'prices: {y_data}')
18  print(f'data_size: {N}')
```

```
areas: [6.7, 4.6, 3.5, 5.5]
prices: [9.1, 5.9, 4.6, 6.7]
data_size: 4
```

❖ House price prediction

❖ One-sample training

1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

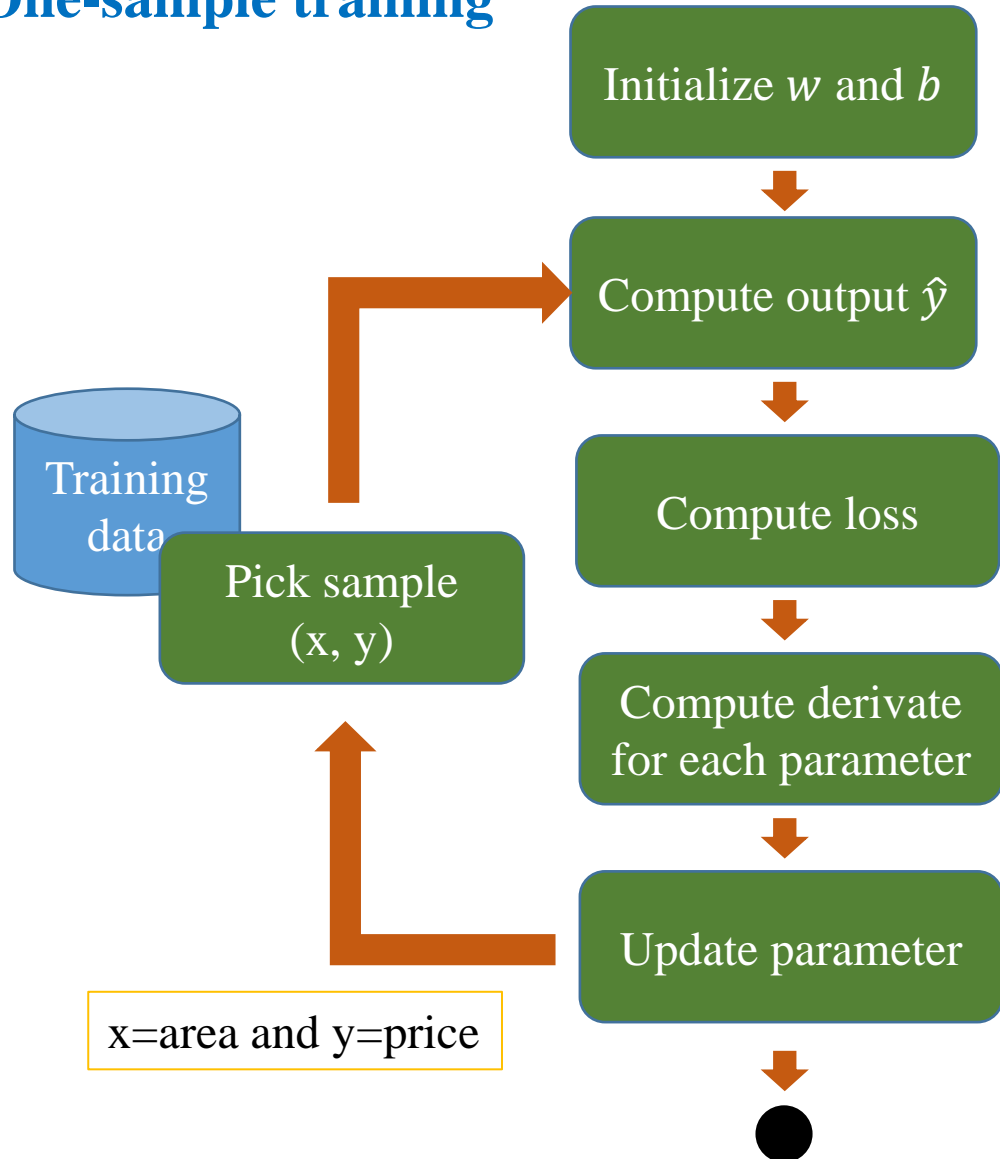
5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w} \quad b = b - \eta \frac{\partial L}{\partial b}$$

```
1 # forward
2 def predict(x, w, b):
3     return x*w + b
4
5 # compute gradient
6 def gradient(y_hat, y, x):
7     dw = 2*x*(y_hat-y)
8     db = 2*(y_hat-y)
9
10    return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17    return (w_new, b_new)
```

Implementation

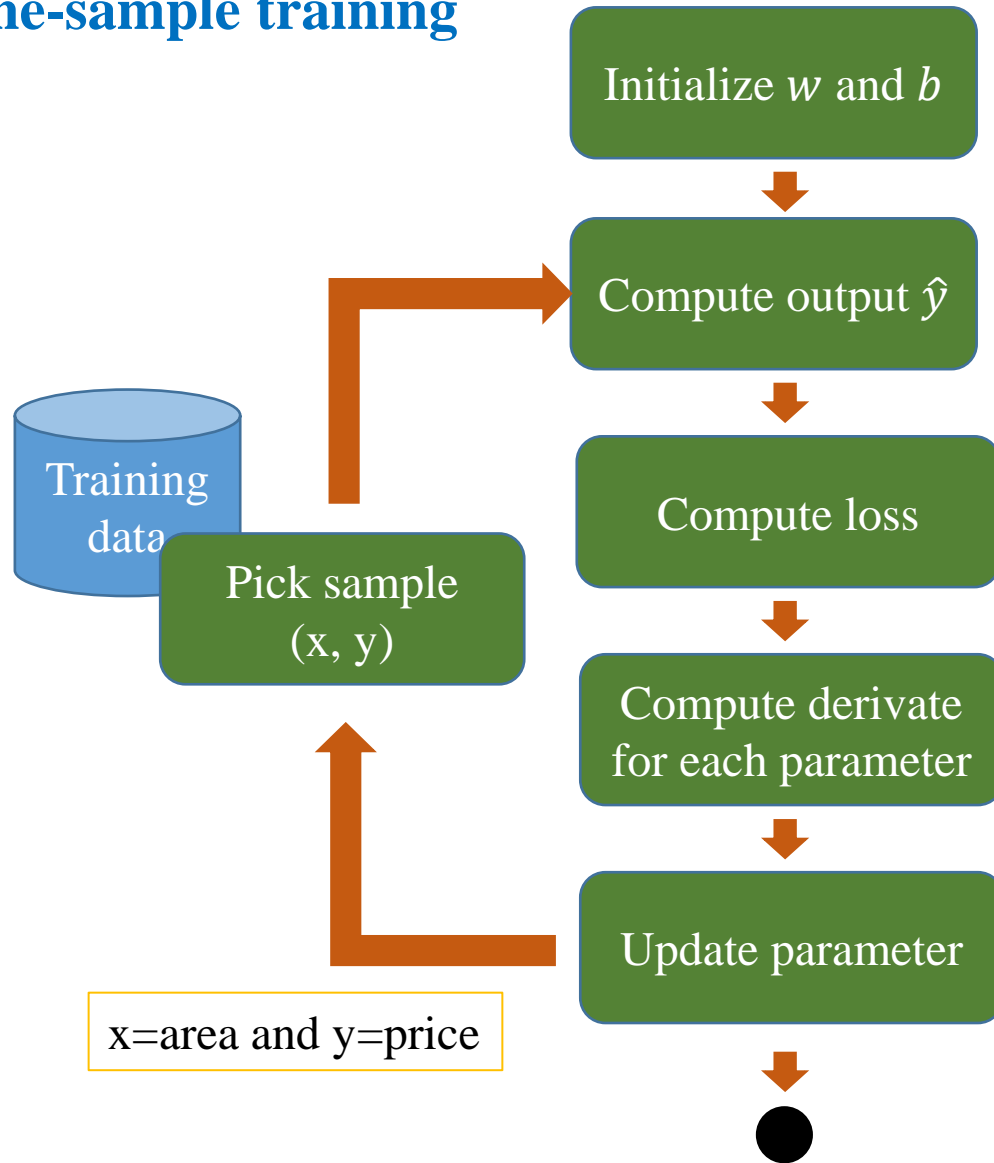
❖ One-sample training



```
1 # init weights
2 b = 0.04
3 w = -0.34
4 lr = 0.01
5
6 # how long
7 epoch_max = 10
8
9 for epoch in range(epoch_max):
10     for i in range(data_size):
11         # get a sample
12         # ...
13
14         # predict y_hat
15         # ...
16
17         # compute loss
18         # ...
19
20         # compute gradient
21         # ...
22
23         # update weights
24         # ...
25
```

Implementation


❖ One-sample training





```
1 # init weights
2 b = 0.04
3 w = -0.34
4 lr = 0.01
5
6 # how long
7 epoch_max = 10
8 data_size = 4
9
10 for epoch in range(epoch_max):
11     for i in range(data_size):
12         # get a sample
13         x = areas[i]
14         y = prices[i]
15
16         # predict y_hat
17         y_hat = predict(x, w, b)
18
19         # compute loss
20         loss = (y_hat-y)*(y_hat-y)
21
22         # compute gradient
23         (dw, db) = gradient(y_hat, y, x)
24
25         # update weights
26         (w, b) = update_weight(w, b, lr, dw, db)
```

Implementation

- ❖ House price prediction
- ❖ One-sample training



Model training



Quiz 1: Construct for the following data

| Features | | Label |
|----------|-------------|---------|
| ↕ Radio | ↕ Newspaper | ↕ Sales |
| 37.8 | 69.2 | 22.1 |
| 39.3 | 45.1 | 10.4 |
| 45.9 | 69.3 | 12 |
| 41.3 | 58.5 | 16.5 |
| 10.8 | 58.4 | 17.9 |
| | | |

Outline

SECTION 1

Linear Regression

SECTION 2

Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions

7000 1250

1) Pick m samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < m$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < m$$

4) Compute derivatives


$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$
$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)})$$

for $0 \leq i < m$

5) Update


$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{m} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{m}$$

❖ Compute derivate for w and b




mini-batch $m = 2$


❖ Compute derivate for w and b




❖ Compute derivate for w and b



❖ Compute derivate for w and b



❖ Compute derivate for w and b

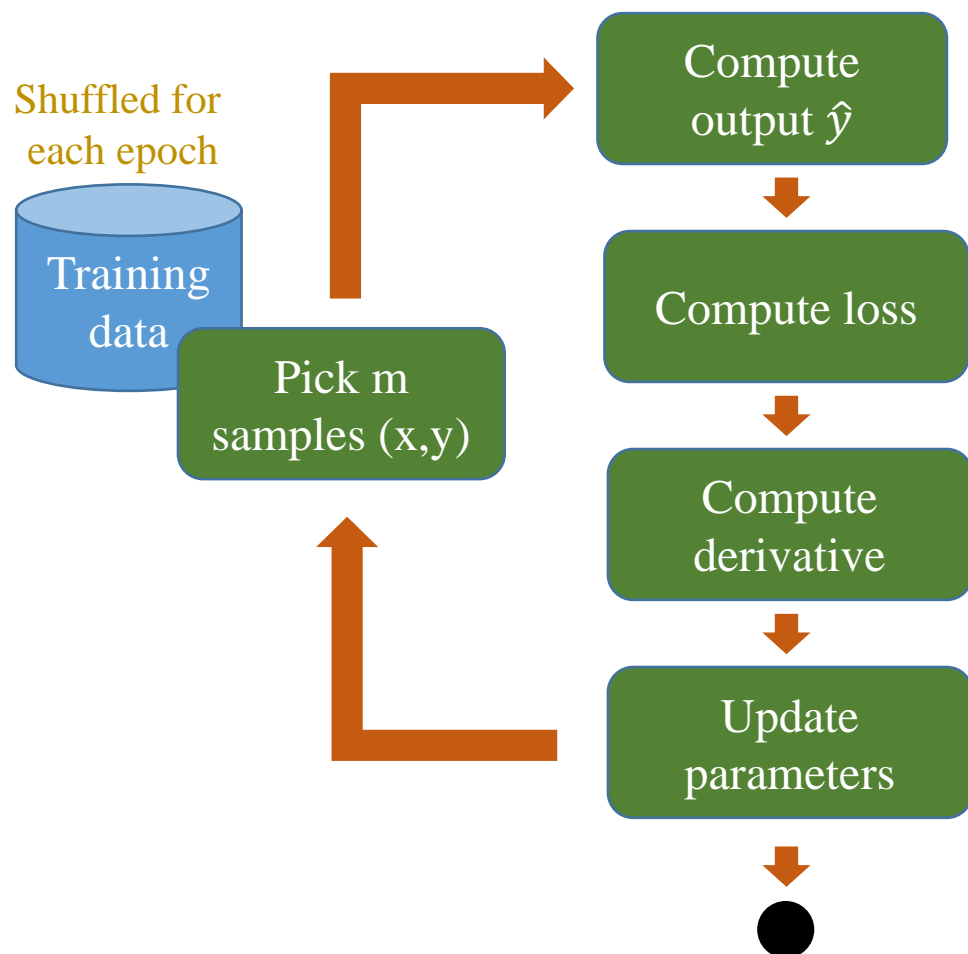


$$\frac{\partial L}{\partial w} = \frac{1}{m} \sum_i 2x^{(i)} (\hat{y}^{(i)} - y^{(i)})$$

$$\frac{\partial L}{\partial b} = \frac{1}{m} \sum_i 2(\hat{y}^{(i)} - y^{(i)})$$

❖ House price prediction

❖ m-sample training ($1 < m < N$)



1) Pick m samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < m$$

3) Compute loss

$$L = \frac{1}{m} \sum_i (\hat{y}^{(i)} - y^{(i)})^2$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$

for $0 \leq i < m$

$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)})$$


5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{m}$$

$$b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{m}$$

❖ House price prediction

❖ m-sample training ($1 < m < N$)




Computational graph

❖ House price prediction

❖ m-sample training ($1 < m < N$)

| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |

$m = 2$



Computational graph

❖ House price prediction

❖ m-sample training ($1 < m < N$)

| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |

$m = 2$

Update w and b

$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$

Learning rate $\eta = 0.01$

replace

replace

$$w = -0.34 - (0.01 * (-110.115)) = 0.761$$

$$b = 0.04 - (0.01 * (-18.762)) = 0.227$$

$$x = \begin{bmatrix} 6.7 \\ 4.6 \end{bmatrix}$$

$$w = -0.34$$

$$\frac{\partial L}{\partial w} = -110.115$$

$$b = 0.04$$

$$\frac{\partial L}{\partial b} = -18.762$$

$$y = \begin{bmatrix} 9.1 \\ 5.9 \end{bmatrix}$$

*

t

+

$(\cdot)^2$


$$t = wx$$

$$\hat{y} = t + b$$


$$L = (\hat{y} - y)^2$$

❖ House price prediction


❖ m-sample training ($1 < m < N$)



Losses for 30 iterations



Model updating for different iterations



Outline

19 5 > 1

SECTION 1

Linear Regression

SECTION 2

Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions

1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < N$$

4) Compute derivatives

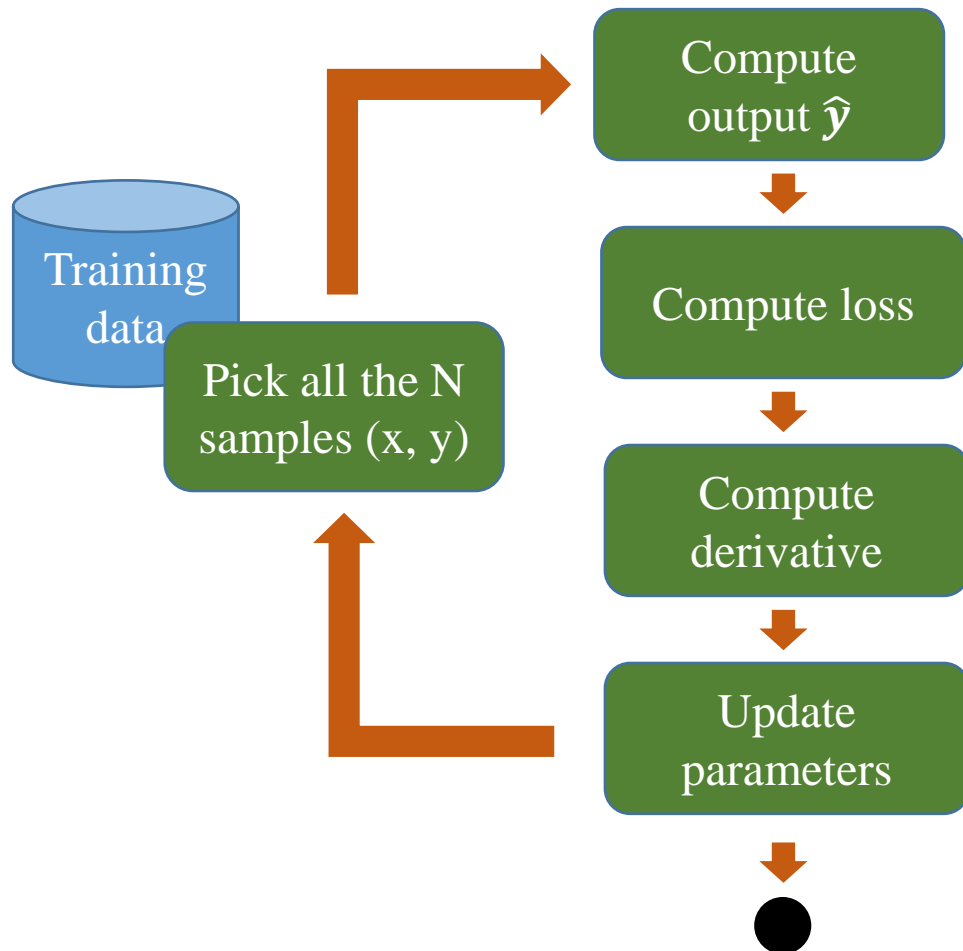
$$\begin{aligned} \frac{\partial L^{(i)}}{\partial w} &= 2x^{(i)}(\hat{y}^{(i)} - y^{(i)}) \\ \frac{\partial L^{(i)}}{\partial b} &= 2(\hat{y}^{(i)} - y^{(i)}) \end{aligned} \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

❖ House price prediction

❖ N-sample training



1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L = \frac{1}{N} \sum_i (\hat{y}^{(i)} - y^{(i)})^2$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$


$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)}) \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L}{\partial b}}{N}$$

❖ House price prediction

❖ N-sample training




Computational graph

❖ House price prediction

❖ N-sample training

$$\frac{\text{sum}(\frac{\partial L}{\partial w})}{4} = -88.5775$$

$$\frac{\text{sum}(\frac{\partial L}{\partial b})}{4} = -16.521$$



❖ House price prediction

❖ N-sample training

Update w and b


$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$

Learning rate $\eta = 0.01$


$$w = -0.34 - (0.01 * (-88.5775)) = 0.54$$

$$b = 0.04 - (0.01 * (-16.521)) = 0.205$$




❖ House price prediction

❖ N-sample training



Losses for 30 iterations



Model updating for different iterations

Extension

| Features | | | Label |
|----------|---------|-------------|---------|
| TV | ↕ Radio | ↕ Newspaper | ↕ Sales |
| 230.1 | 37.8 | 69.2 | 22.1 |
| 44.5 | 39.3 | 45.1 | 10.4 |
| 17.2 | 45.9 | 69.3 | 12 |
| 151.5 | 41.3 | 58.5 | 16.5 |
| 180.8 | 10.8 | 58.4 | 17.9 |

Advertising data

Model: $\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$

$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$

❖ General formula

| Feature | Label |
|---------|-------|
| area | price |
| 6.7 | 9.1 |
| 4.6 | 5.9 |
| 3.5 | 4.6 |
| 5.5 | 6.7 |

House price data

Model: $\hat{y} = w_1x_1 + b$

price = a * area + b

| Features | | | Label |
|----------|---------|-------------|---------|
| TV | ↕ Radio | ↕ Newspaper | ↕ Sales |
| 230.1 | 37.8 | 69.2 | 22.1 |
| 44.5 | 39.3 | 45.1 | 10.4 |
| 17.2 | 45.9 | 69.3 | 12 |
| 151.5 | 41.3 | 58.5 | 16.5 |
| 180.8 | 10.8 | 58.4 | 17.9 |

Advertising data

Model: $\hat{y} = \underline{w_1}x_1 + \underline{w_2}x_2 + \underline{w_3}x_3 + b$

Sale = $w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$

Linear Regression

1) Pick a sample (x_1, x_2, x_3, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = w_1 * TV + w_2 * R + w_3 * N + b$$

$$\hat{y} = w_1 * x_1 + w_2 * x_2 + w_3 * x_3 + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w_1} = 2x_1(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_3} = 2x_3(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_2} = 2x_2(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w_1 = w_1 - \eta \frac{\partial L}{\partial w_1}$$

$$w_3 = w_3 - \eta \frac{\partial L}{\partial w_3}$$

$$w_2 = w_2 - \eta \frac{\partial L}{\partial w_2}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

| Features | | | Label |
|----------|-------|-----------|-------|
| TV | Radio | Newspaper | Sales |
| 230.1 | 37.8 | 69.2 | 22.1 |
| 44.5 | 39.3 | 45.1 | 10.4 |
| 17.2 | 45.9 | 69.3 | 12 |
| 151.5 | 41.3 | 58.5 | 16.5 |
| 180.8 | 10.8 | 58.4 | 17.9 |

Advertising data

Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$

1) Pick a sample (x_1, x_2, x_3, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = w_1 * TV + w_2 * R + w_3 * N + b$$

$$\hat{y} = w_1 * x_1 + w_2 * x_2 + w_3 * x_3 + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w_1} = 2x_1(\hat{y} - y) \quad \frac{\partial L}{\partial w_3} = 2x_3(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_2} = 2x_2(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w_1 = w_1 - \eta \frac{\partial L}{\partial w_1} \quad w_3 = w_3 - \eta \frac{\partial L}{\partial w_3}$$

$$w_2 = w_2 - \eta \frac{\partial L}{\partial w_2} \quad b = b - \eta \frac{\partial L}{\partial b}$$

```
1  # compute output and loss
2  def predict(x1, x2, x3, w1, w2, w3, b):
3      return w1*x1 + w2*x2 + w3*x3 + b
4  def compute_loss(y_hat, y):
5      return (y_hat - y)**2
6
7  # compute gradient
8  def compute_gradient_wi(xi, y, y_hat):
9      dl_dwi = 2*xi*(y_hat-y)
10     return dl_dwi
11 def compute_gradient_b(y, y_hat):
12     dl_db = 2*(y_hat-y)
13     return dl_db
14
15 # update weights
16 def update_weight_wi(wi, dl_dwi, lr):
17     wi = wi - lr*dl_dwi
18     return wi
19 def update_weight_b(b, dl_db, lr):
20     b = b - lr*dl_db
21     return b
```

| Features | | | Label |
|----------|---------|-------------|---------|
| TV | ↕ Radio | ↕ Newspaper | ↕ Sales |
| 230.1 | 37.8 | 69.2 | 22.1 |
| 44.5 | 39.3 | 45.1 | 10.4 |
| 17.2 | 45.9 | 69.3 | 12 |
| 151.5 | 41.3 | 58.5 | 16.5 |
| 180.8 | 10.8 | 58.4 | 17.9 |

```

1 def initialize_params():
2     w1 = random.gauss(mu=0.0, sigma=0.01)
3     w2 = random.gauss(mu=0.0, sigma=0.01)
4     w3 = random.gauss(mu=0.0, sigma=0.01)
5     b = 0
6
7     return w1, w2, w3, b
8
9 # initialize model's parameters
10 w1, w2, w3, b = initialize_params()
11 print(w1, w2, w3, b)

```

0.01609506469549467 0.00607778501208891 0.0023344573891806507 0

```

1 import numpy as np
2 import random
3
4 def get_column(data, index):
5     result = [row[index] for row in data]
6     return result
7
8 data = np.genfromtxt('advertising.csv',
9                     delimiter=',',
10                    skip_header=1).tolist()
11
12 # get tv (index=0)
13 tv_data = get_column(data, 0)
14
15 # get radio (index=1)
16 radio_data = get_column(data, 1)
17
18 # get newspaper (index=2)
19 newspaper_data = get_column(data, 2)
20
21 # get sales (index=3)
22 sales_data = get_column(data, 3)

```

Outline

SECTION 1

Linear Regression

SECTION 2


Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions (Optional)



Discussion 1: Is it OK to use the following loss function?

$$L = \frac{1}{2}(\hat{y} - y)^2$$

Discussion 2: if so, construct formulas

$$\frac{\partial L}{\partial y} = (\hat{y} - y) \cdot 1$$

Discussion 3: What about the following loss function?

$$L = \frac{1}{2} (y - \hat{y})^2$$

$(y - \hat{y})^2$

$\frac{1}{2}$

Discussion 4: Construct the connection between the two following losses?


$$L_1 = \frac{1}{2}(\hat{y} - y)^2$$

$$L_2 = (\hat{y} - y)^2$$

$$\hat{y} - y$$

$$\hat{y} = wx + b$$

Discussion 5:
Can we remove b?



$$\hat{y} = wx + b$$

Discussion 5:
Can we remove b?

