



Module 04 – Exercise Class

LINEAR REGRESSION


Vectorization

Nguyen Quoc Thai

Objectives

Linear Regression

- ❖ Introduction
- ❖ Stochastic Gradient Descent
- ❖ Mini Batch Gradient Descent
- ❖ Batch Gradient Descent




Linear Regression for Timeseries

- ❖ Bitcoin Dataset
- ❖ Feature Scaling
- ❖ Modeling
- ❖ Evaluation
- ❖ Inference

Outline

SECTION 1

Linear Regression



SECTION 2

Time Series Application



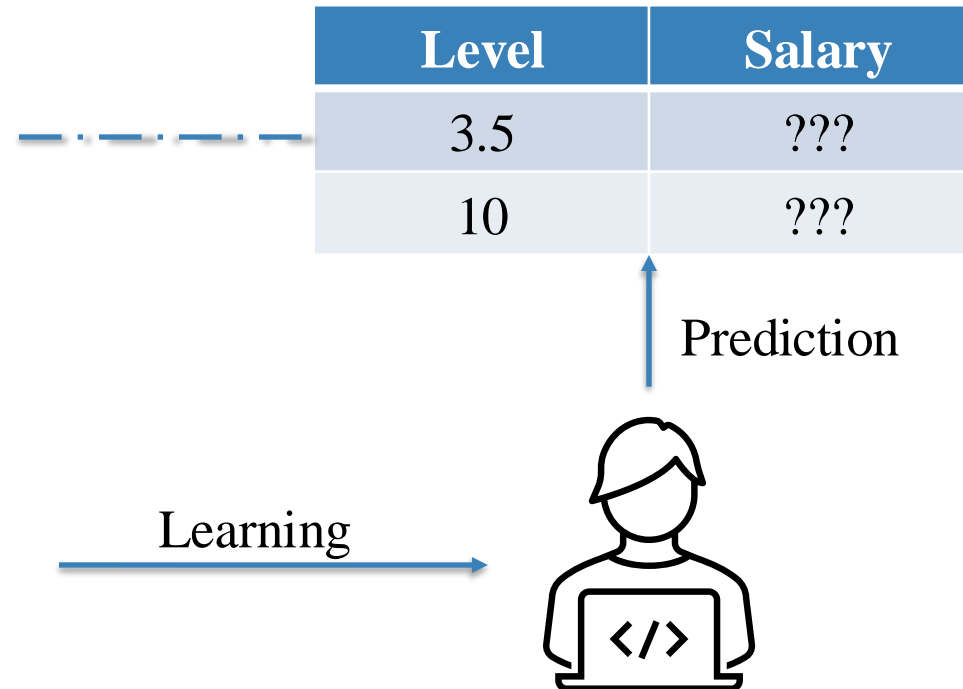
Linear Regression



Introduction

Data

Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47



Linear Regression




Introduction

Data

Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47

Visualization



Linear Regression



Introduction

Data

Level	Salary
0	8
1	15
2	18
3	22
4	26
5	30
6	38
7	47


Modeling

$$y = wx + b$$



Find w and b to fit the data

Visualization



Linear Regression



Stochastic Gradient Descent

1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$


$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate



Difference between
predicted and actual value

Linear Regression

! Stochastic Gradient Descent (Vectorization)

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \boldsymbol{\theta}^T \mathbf{x} = \mathbf{x}^T \boldsymbol{\theta}$$

3) Compute loss

$$L = (\hat{y} - y)^2$$


4) Compute derivative

$$\nabla_{\boldsymbol{\theta}} L = 2\mathbf{x}(\hat{y} - y)$$

5) Update parameters

$$\boldsymbol{\theta} = \boldsymbol{\theta} - \eta \nabla_{\boldsymbol{\theta}} L$$

η is learning rate



Difference between
predicted and actual value

Linear Regression

! Stochastic Gradient Descent (Vectorization)

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
x = np.array([[2]])  
y = np.array([18])
```

```
num_samples = x.shape[0]  
num_samples
```

```
1
```

```
# append bias  
x = np.hstack([x, np.ones((num_samples, 1))])  
x
```

```
array([[2., 1.]])
```

```
# init weights  
theta = np.array([2, 2]) # num_features: x.shape[1]  
theta
```

```
array([2, 2])
```

Linear Regression

! Stochastic Gradient Descent (Vectorization)

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
# forward
def predict(x, theta):
    return x.dot(theta)

y_hat = predict(x, theta)
y_hat

array([6.])
```

Linear Regression

! Stochastic Gradient Descent (Vectorization)

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
# compute gradient
```

```
def compute_gradient(x, y, y_hat):
```

```
    d_theta = 2*x*(y_hat-y)
```

```
    return d_theta
```

```
d_theta = compute_gradient(x, y, y_hat)
```

```
d_theta
```

```
array([[ -48.,  -24.]])
```

Linear Regression

! Stochastic Gradient Descent (Vectorization)

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
# update weights
```

```
lr = 0.1
```

```
def update_weights(theta, d_theta, lr):
```

```
    new_theta = theta - lr*d_theta
```

```
    return new_theta
```

```
new_theta = update_weights(theta, d_theta, lr)
```

```
new_theta
```

```
array([[6.8, 4.4]])
```

Linear Regression



Practice

1) Pick a sample (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

$$x = [1 \ 3]$$
$$y = 22$$

Init θ
lr = 0.1

$$y = 2x + 2$$

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 x = np.array([[2], [3], [1]])
2 y = np.array([3, 4, 2])
3 x, y
```

✓ 0.0s

```
(array([[2],
        [3],
        [1]]),
 array([3, 4, 2]))
```

```
1 x.shape, y.shape
```

✓ 0.0s

```
((3, 1), (3,))
```

```
1 num_samples = x.shape[0]
```

```
2 num_samples
```

✓ 0.0s

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 # append bias
2 x = np.hstack([np.ones((num_samples, 1)), x])
3 x
✓ 0.0s
```

```
array([[1., 2.],
       [1., 3.],
       [1., 1.]])
```

```
1 # init weights
2 theta = np.array([2, 2]) # num_features: x.shape[1]
3 theta
✓ 0.0s
```

```
array([2, 2])
```

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 # forward
2 def predict(x, theta):
3     return x.dot(theta)
4
5 y_hat = predict(x, theta)
6 y_hat
```

✓ 0.0s

array([6., 8., 4.])

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 # compute loss
2 def compute_loss(y_hat, y):
3     return (y_hat-y)*(y_hat-y)
4
5 loss = compute_loss(y_hat, y)
6 loss
```

✓ 0.0s

array([9., 16., 4.])

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 # compute gradient
2 def compute_gradient(x, y, y_hat):
3     d_theta = 2*x.T.dot(y_hat-y)
4     return d_theta
5
6 d_theta = compute_gradient(x, y, y_hat)
7 d_theta
```

✓ 0.0s

array([18., 40.])

Linear Regression



Mini Batch Gradient Descent (Vectorization)

1) Pick m samples (x, y) from training data

2) Compute output \hat{y}

$$\hat{y} = \theta^T x = x^T \theta$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\nabla_{\theta} L = 2x(\hat{y} - y)$$

5) Update parameters

$$\theta = \theta - \eta \nabla_{\theta} L$$

η is learning rate

```
1 # update weights
2 lr = 0.1
3
4 def update_weights(theta, d_theta, lr):
5     new_theta = theta - lr*d_theta
6     return new_theta
7
8 new_theta = update_weights(theta, d_theta, lr)
9 new_theta
```


✓ 0.0s

array([0.2, -2.])

Linear Regression



Comparison




QUIZ TIME

Outline

SECTION 1

Linear Regression



SECTION 2

Time Series Prediction



Time Series Prediction



Bitcoin Dataset

```
1 df = pd.read_csv('/content/BTC-Daily.csv')
2 df = df.drop_duplicates()
3
```

```
1 df.head()
```

	unix	date	symbol	open	high	low	close	Volume BTC	Volume USD
0	1646092800	2022-03-01 00:00:00	BTC/USD	43221.71	43626.49	43185.48	43185.48	49.006289	2.116360e+06
1	1646006400	2022-02-28 00:00:00	BTC/USD	37717.10	44256.08	37468.99	43178.98	3160.618070	1.364723e+08
2	1645920000	2022-02-27 00:00:00	BTC/USD	39146.66	39886.92	37015.74	37712.68	1701.817043	6.418008e+07
3	1645833600	2022-02-26 00:00:00	BTC/USD	39242.64	40330.99	38600.00	39146.66	912.724087	3.573010e+07
4	1645747200	2022-02-25 00:00:00	BTC/USD	38360.93	39727.97	38027.61	39231.64	2202.851827	8.642149e+07

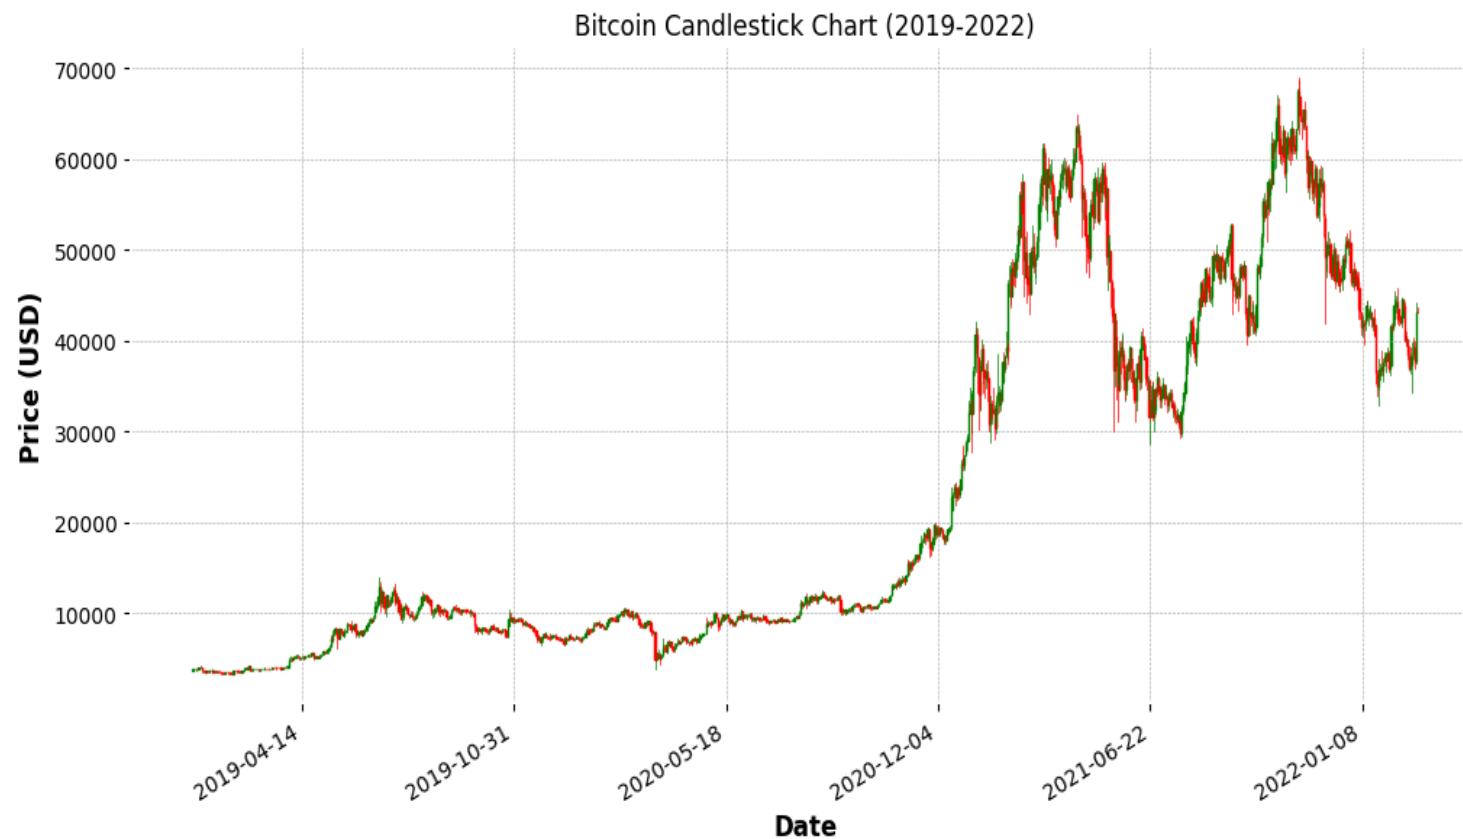
```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2651 entries, 0 to 2650
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   unix            2651 non-null   int64
1   date            2651 non-null   object
2   symbol          2651 non-null   object
3   open            2651 non-null   float64
4   high            2651 non-null   float64
5   low             2651 non-null   float64
6   close           2651 non-null   float64
7   Volume BTC      2651 non-null   float64
8   Volume USD      2651 non-null   float64
dtypes: float64(6), int64(1), object(2)
memory usage: 186.5+ KB
```

Time Series Prediction



Bitcoin Dataset






Time Series Prediction



Bitcoin Dataset



Time Series Prediction



Feature Scaling

MaxAbsScaler

$$x_{new} = \frac{x}{x_{max}}$$

1	0.25
2	0.5
3	0.75
4	1.0

```
1 import numpy as np
2 from sklearn.preprocessing import MaxAbsScaler
3
4 X = np.array([[1], [2], [3], [4]])
5
6 scaler = MaxAbsScaler()
7 X_transform = scaler.fit_transform(X)
8 X_transform
```

✓ 0.0s

```
array([[0.25],
       [0.5 ],
       [0.75],
       [1.  ]])
```

Time Series Prediction



Feature Scaling

MinMaxScaler

$$x_{new} = \frac{x - x_{min}}{x_{max} - x_{min}}$$

1	0.0
2	0.33
3	0.67
4	1.0

```
1 import numpy as np
2 from sklearn.preprocessing import MinMaxScaler
3
4 X = np.array([[1], [2], [3], [4]])
5
6 scaler = MinMaxScaler()
7 X_transform = scaler.fit_transform(X)
8 X_transform
```

✓ 0.0s

```
array([[0.        ],
       [0.33333333],
       [0.66666667],
       [1.        ]])
```

Time Series Prediction



Feature Scaling

StandardScaler

$$x_{new} = \frac{x - \mu}{\sigma}$$

1	-1.34
2	-0.45
3	0.45
4	1.34

$$\mu = \frac{1}{N} \sum x_i$$

$$\sigma = \sqrt{\frac{1}{N} \sum_i (x_i - \mu)^2}$$

```
1 import numpy as np
2 from sklearn.preprocessing import StandardScaler
3
4 X = np.array([[1], [2], [3], [4]])
5
6 scaler = StandardScaler()
7 X_transform = scaler.fit_transform(X)
8 X_transform
```

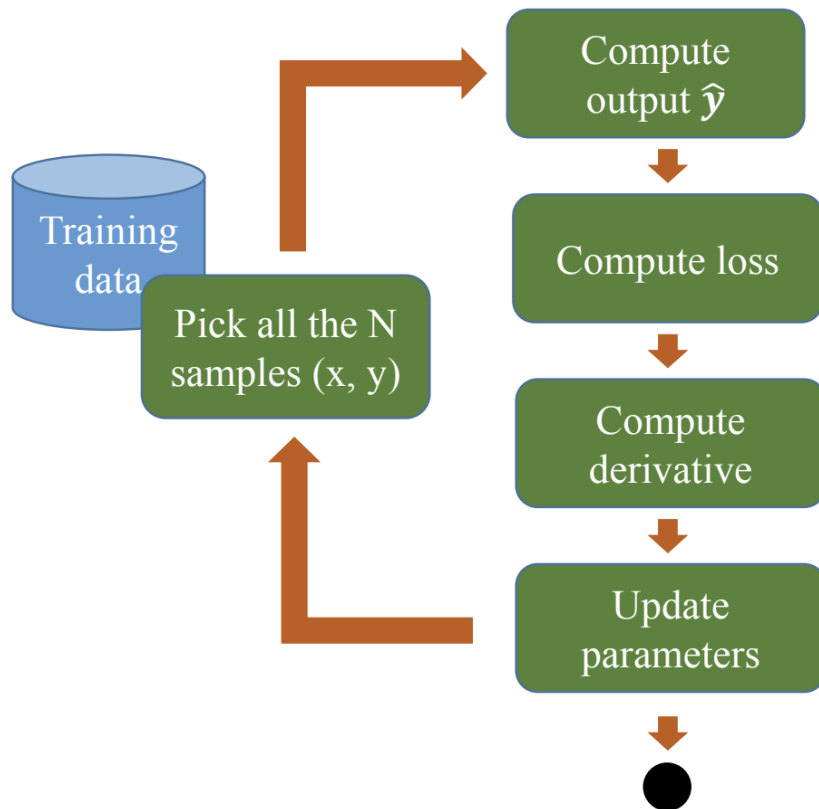
✓ 0.0s

```
array([[ -1.34164079],
       [ -0.4472136 ],
       [  0.4472136 ],
       [  1.34164079]])
```

Time Series Prediction



Modeling



1) Pick all the N samples from training data

2) Compute output \hat{y}

$$\hat{y} = x\theta$$

3) Compute loss

$$L(\hat{y}, y) = (\hat{y} - y) \odot (\hat{y} - y)$$

4) Compute derivative

$$k = 2(\hat{y} - y)$$

$$L'_{\theta} = x^T k$$


5) Update parameters

$$\theta = \theta - \eta \frac{L'_{\theta}}{N} \quad \eta \text{ is learning rate}$$

Summary

Linear Regression

- ❖ Introduction
- ❖ Stochastic Gradient Descent
- ❖ Mini Batch Gradient Descent
- ❖ Batch Gradient Descent



Linear Regression for Timeseries

- ❖ Bitcoin Dataset
- ❖ Feature Scaling
- ❖ Modeling
- ❖ Evaluation
- ❖ Inference



AI VIET NAM

@aivietnam.edu.vn

Thanks!

Any questions?